

NLP Preprocessing And Text Classification

Course Name: MDM Deep Learning

Lab Title: NLP Techniques for Text Classification

Student Name: Rohit Dahale

Student ID: 202201070052

Date of Submission: 16/4/2025

Group Members: Sakshi Aher, Akhilesh Ukey

Objective The objective of this assignment is to implement NLP preprocessing techniques and build a text classification model using machine learning techniques.

Learning Outcomes:

1. Understand and apply NLP preprocessing techniques such as tokenization, stopwords removal, stemming, and lemmatization.
2. Implement text vectorization techniques such as TF-IDF and CountVectorizer.
3. Develop a text classification model using a machine learning algorithm.
4. Evaluate the performance of the model using suitable metrics.

```
from google.colab import files
uploaded = files.upload()

<IPython.core.display.HTML object>

Saving all-data.csv to all-data (1).csv

import pandas as pd

# Load the CSV with proper encoding
df = pd.read_csv("all-data.csv", encoding='ISO-8859-1', header=None)

# Rename columns
df.columns = ['label', 'text']
df.head()

{"summary": "{\n  \"name\": \"df\",\n  \"rows\": 4846,\n  \"fields\": [\n    {\n      \"column\": \"label\",\n      \"properties\": {\n        \"dtype\": \"category\",\n        \"num_unique_values\": 3,\n        \"samples\": [\n          \"neutral\",\n          \"negative\",\n          \"positive\"],\n        \"semantic_type\": \"\",\n        \"description\": \"\"}\n    },\n    {\n      \"column\":
```

```

{"text","\n      \"properties\": {\n          \"dtype\": \"string\",\n          \"num_unique_values\": 4838,\n          \"samples\": [\n              \"The Company serves approximately 3,000 customers in over 100 countries .\",\n              \"On Dec. 1 , Grimaldi acquired 1.5 million shares and a 50.1-percent stake in Finnlines .\",\n              \"The extracted filtrates are very high in clarity while the dried filter cakes meet required transport moisture limits (TMLs)for their ore grades .\"\n          ],\n          \"semantic_type\": \"\",\n          \"description\": \"\"\n      }\n    },\n    \"type\": \"dataframe\", \"variable_name\": \"df\"}

```

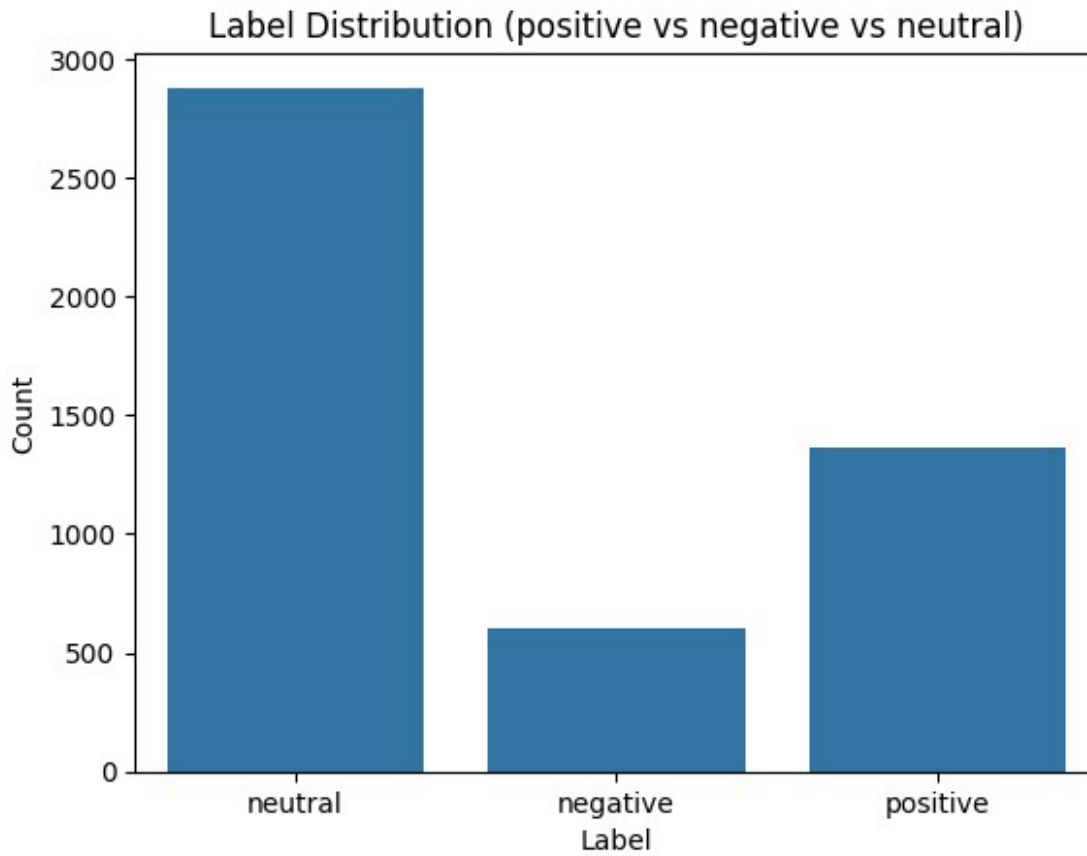
```
# Check for missing values in any column
print("Missing values:\n", df.isnull().sum())
```

```
# Print data types of each column
print("\nData Types:\n", df.dtypes)
```

```
# Visualize the distribution of labels (ham vs spam)
sns.countplot(x='label', data=df)
plt.title("Label Distribution (positive vs negative vs neutral)")
plt.xlabel("Label")
plt.ylabel("Count")
plt.show()
```

```
Missing values:
  label      0
text        0
processed_text  0
dtype: int64
```

```
Data Types:
  label      object
text         object
processed_text  object
dtype: object
```



```
import nltk

nltk.download('punkt')           # For tokenization
nltk.download('stopwords')       # For removing stopwords
nltk.download('wordnet')         # For lemmatization
nltk.download('omw-1.4')         # WordNet data
nltk.download('punkt_tab')       # <- Specific one mentioned in error

[nltk_data] Downloading package punkt to /root/nltk_data...
[nltk_data]   Package punkt is already up-to-date!
[nltk_data] Downloading package stopwords to /root/nltk_data...
[nltk_data]   Package stopwords is already up-to-date!
[nltk_data] Downloading package wordnet to /root/nltk_data...
[nltk_data]   Package wordnet is already up-to-date!
[nltk_data] Downloading package omw-1.4 to /root/nltk_data...
[nltk_data]   Package omw-1.4 is already up-to-date!
[nltk_data] Downloading package punkt_tab to /root/nltk_data...
[nltk_data]   Package punkt_tab is already up-to-date!

True

from nltk.tokenize import word_tokenize
from nltk.corpus import stopwords
```

```

from nltk.stem import PorterStemmer, WordNetLemmatizer
import string

stop_words = set(stopwords.words('english'))
stemmer = PorterStemmer()
lemmatizer = WordNetLemmatizer()

def preprocess_text(text):
    text = text.lower()
    text = text.translate(str.maketrans('', '', string.punctuation))
    tokens = word_tokenize(text)
    tokens = [w for w in tokens if w not in stop_words]
    tokens = [stemmer.stem(lemmatizer.lemmatize(w)) for w in tokens]
    return ' '.join(tokens)

df['processed_text'] = df['text'].apply(preprocess_text)
df[['label', 'processed_text']].head()

{"summary": "{\n  \"name\": \"df[['label', 'processed_text']]\", \n  \"rows\": 5, \n  \"fields\": [\n    {\n      \"column\": \"label\", \n      \"properties\": {\n        \"dtype\": \"string\", \n        \"num_unique_values\": 3, \n        \"samples\": [\n          \"neutral\", \n          \"negative\", \n          \"positive\" \n        ], \n        \"semantic_type\": \"\", \n        \"description\": \"\" \n      } \n    }, \n    {\n      \"column\": \"processed_text\", \n      \"properties\": {\n        \"dtype\": \"string\", \n        \"num_unique_values\": 5, \n        \"samples\": [\n          \"technopoli plan develop stage area less 100000 squar meter order\n          host compani work comput technolog telecommun statement said\", \n          \"accord compani updat strategi year 20092012 baswar target longterm\n          net sale growth rang 20 40 oper profit margin 10 20 net sale\", \n          \"intern electron industri compani elcoteq laid ten employe tallinn\n          facil contrari earlier layoff compani contract rank offic worker daili\n          postime report\" \n        ], \n        \"semantic_type\": \"\", \n        \"description\": \"\" \n      } \n    ] \n  }\", \"type\": \"dataframe\"}

from sklearn.feature_extraction.text import CountVectorizer,
TfidfVectorizer

# TF-IDF Vectorizer
tfidf = TfidfVectorizer()
X_tfidf = tfidf.fit_transform(df['processed_text'])

# Count Vectorizer
count_vec = CountVectorizer()
X_count = count_vec.fit_transform(df['processed_text'])

# Target labels
y = df['label']

```

```

from sklearn.model_selection import train_test_split

# Using TF-IDF for modeling
X_train, X_test, y_train, y_test = train_test_split(X_tfidf, y,
test_size=0.2, random_state=42)

from sklearn.linear_model import LogisticRegression

model = LogisticRegression(max_iter=1000)
model.fit(X_train, y_train)
y_pred = model.predict(X_test)

from sklearn.metrics import accuracy_score, precision_score,
recall_score, f1_score, classification_report, confusion_matrix
import seaborn as sns
import matplotlib.pyplot as plt

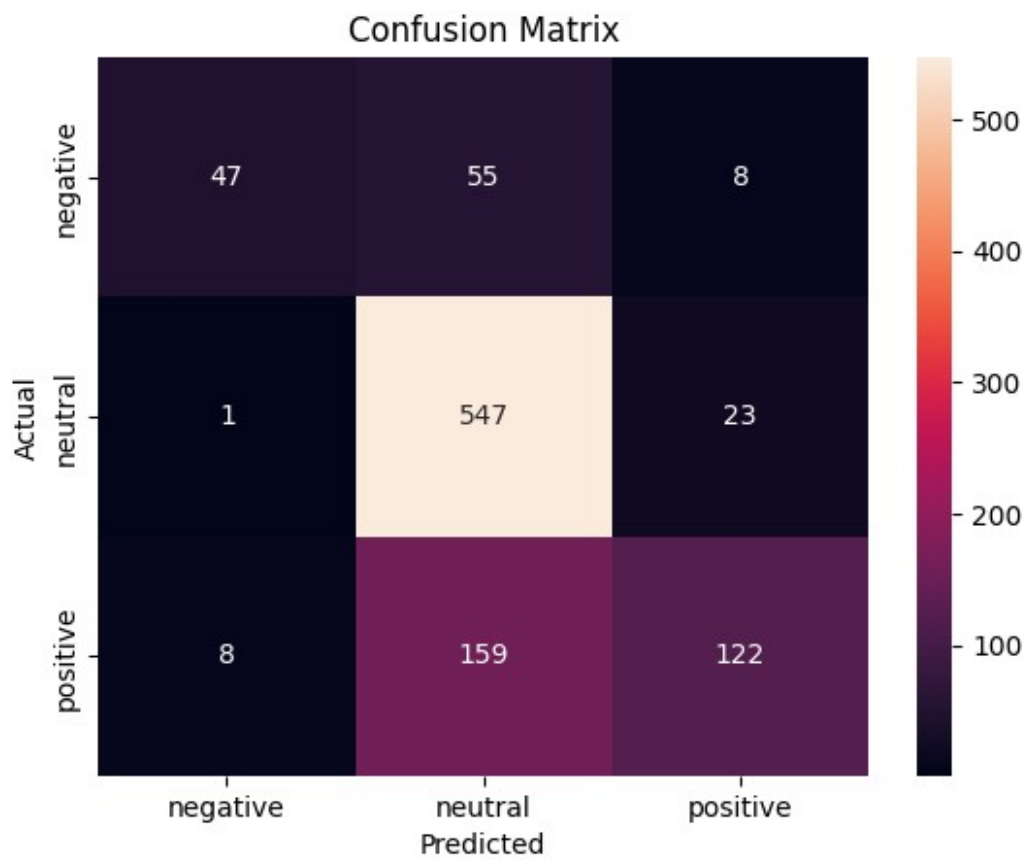
print("Accuracy:", accuracy_score(y_test, y_pred))
print("Precision:", precision_score(y_test, y_pred,
average='weighted'))
print("Recall:", recall_score(y_test, y_pred, average='weighted'))
print("F1 Score:", f1_score(y_test, y_pred, average='weighted'))

# Confusion Matrix
conf_mat = confusion_matrix(y_test, y_pred, labels=model.classes_)
sns.heatmap(conf_mat, annot=True, fmt='d', xticklabels=model.classes_,
yticklabels=model.classes_)
plt.xlabel("Predicted")
plt.ylabel("Actual")
plt.title("Confusion Matrix")
plt.show()

# Classification Report
print("\nClassification Report:\n")
print(classification_report(y_test, y_pred))

Accuracy: 0.7381443298969073
Precision: 0.7558717224556394
Recall: 0.7381443298969073
F1 Score: 0.7121671058138577

```



Classification Report:

	precision	recall	f1-score	support
negative	0.84	0.43	0.57	110
neutral	0.72	0.96	0.82	571
positive	0.80	0.42	0.55	289
accuracy			0.74	970
macro avg	0.79	0.60	0.65	970
weighted avg	0.76	0.74	0.71	970

Discussion and Conclusion :

After implementing the text classification pipeline using natural language processing (NLP) techniques on the Financial news dataset, the following model evaluation metrics were observed:

Accuracy: 0.73

Precision: 0.75

Recall: 0.73
F1 Score: 0.71

Declaration

I, Rohit Dahale, confirm that the work submitted in this assignment is my own and has been completed following academic integrity guidelines. The code is uploaded on my GitHub repository account, and the repository link is provided below:

GitHub Repository Link: <https://github.com/rohitdahale/NLP-Techniques>

Signature: Rohit R. Dahale