

Course Name: Deep Learning

Lab Title: Experiment 5.3: Sequence Text Classification using LSTM

[Dataset](#)

Name - Rohit Dahale **PRN** - 202201070052

Objective: To classify text sequences using LSTM-based models (e.g., sentiment or spam detection).

```
!pip install numpy pandas matplotlib seaborn scikit-learn tensorflow --quiet
```

```
#Import libraries
import pandas as pd
import numpy as np
import seaborn as sns
import matplotlib.pyplot as plt
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import LabelEncoder
from sklearn.metrics import accuracy_score, precision_score, f1_score, confusion_matrix, classification_report
```

```
import tensorflow as tf
from tensorflow.keras.preprocessing.text import Tokenizer
from tensorflow.keras.preprocessing.sequence import pad_sequences
from tensorflow.keras.models import Sequential
from tensorflow.keras.layers import Embedding, LSTM, Dense, Dropout, Bidirectional
```

```
#Load your spam dataset
df = pd.read_csv("/content/spam.csv", encoding='ISO-8859-1')
df = df[['v1', 'v2']]
df.columns = ['label', 'message']
df.head()
```



	label	message
0	ham	Go until jurong point, crazy.. Available only ...
1	ham	OK lar... Joking wif u oni...
2	spam	Free entry in 2 a wkly comp to win FA Cup fina...
3	ham	U dun say so early hor... U c already then say...
4	ham	Nah I don't think he goes to usf, he lives aro...

```
#Encode labels and split
le = LabelEncoder()
df['label_num'] = le.fit_transform(df['label'])

X_train, X_test, y_train, y_test = train_test_split(df['message'], df['label_num'], test_size=0.2, random_state=42)
```

```
max_words = 5000
max_len = 100
```

```
tokenizer = Tokenizer(num_words=max_words, oov_token='<OOV>')
tokenizer.fit_on_texts(X_train)
```

```
X_train_seq = tokenizer.texts_to_sequences(X_train)
X_test_seq = tokenizer.texts_to_sequences(X_test)
```

```
X_train_pad = pad_sequences(X_train_seq, maxlen=max_len, padding='post')
X_test_pad = pad_sequences(X_test_seq, maxlen=max_len, padding='post')
```

```
# Build Simple LSTM
model1 = Sequential([
    Embedding(max_words, 64, input_length=max_len),
    LSTM(64),
    Dense(1, activation='sigmoid')
])
model1.compile(loss='binary_crossentropy', optimizer='adam', metrics=['accuracy'])
model1.summary()
```

```
# Train
model1.fit(X_train_pad, y_train, epochs=5, batch_size=32, validation_split=0.1)
```

```

/usr/local/lib/python3.11/dist-packages/keras/src/layers/core/embedding.py:90: UserWarning: Argument `input_length` is deprecated.
warnings.warn(
Model: "sequential"

```

Layer (type)	Output Shape	Param #
embedding (Embedding)	?	0 (unbuilt)
lstm (LSTM)	?	0 (unbuilt)
dense (Dense)	?	0 (unbuilt)

Total params: 0 (0.00 B)

Trainable params: 0 (0.00 B)

Non-trainable params: 0 (0.00 B)

Epoch 1/5

126/126 ————— 12s 61ms/step - accuracy: 0.8688 - loss: 0.4389 - val_accuracy: 0.8565 - val_loss: 0.4127

Epoch 2/5

126/126 ————— 9s 75ms/step - accuracy: 0.8620 - loss: 0.4032 - val_accuracy: 0.8565 - val_loss: 0.4131

Epoch 3/5

126/126 ————— 10s 74ms/step - accuracy: 0.8555 - loss: 0.4143 - val_accuracy: 0.8565 - val_loss: 0.4167

Epoch 4/5

126/126 ————— 9s 72ms/step - accuracy: 0.8664 - loss: 0.3948 - val_accuracy: 0.8565 - val_loss: 0.4117

Epoch 5/5

126/126 ————— 9s 61ms/step - accuracv: 0.8706 - loss: 0.3859 - val accuracv: 0.8565 - val loss: 0.4119

BiLSTM Model

```

model2 = Sequential([
    Embedding(max_words, 64, input_length=max_len),
    Bidirectional(LSTM(64)),
    Dense(1, activation='sigmoid')
])
model2.compile(loss='binary_crossentropy', optimizer='adam', metrics=['accuracy'])
model2.summary()

```

Train

```

model2.fit(X_train_pad, y_train, epochs=5, batch_size=32, validation_split=0.1)

```

```

Model: "sequential_1"

```

Layer (type)	Output Shape	Param #
embedding_1 (Embedding)	?	0 (unbuilt)
bidirectional (Bidirectional)	?	0 (unbuilt)
dense_1 (Dense)	?	0 (unbuilt)

Total params: 0 (0.00 B)

Trainable params: 0 (0.00 B)

Non-trainable params: 0 (0.00 B)

Epoch 1/5

126/126 ————— 22s 135ms/step - accuracy: 0.8797 - loss: 0.3529 - val_accuracy: 0.9664 - val_loss: 0.1013

Epoch 2/5

126/126 ————— 18s 114ms/step - accuracy: 0.9836 - loss: 0.0536 - val_accuracy: 0.9776 - val_loss: 0.0788

Epoch 3/5

126/126 ————— 21s 117ms/step - accuracy: 0.9958 - loss: 0.0158 - val_accuracy: 0.9731 - val_loss: 0.0936

Epoch 4/5

126/126 ————— 20s 115ms/step - accuracy: 0.9964 - loss: 0.0150 - val_accuracy: 0.9821 - val_loss: 0.0896

Epoch 5/5

126/126 ————— 15s 115ms/step - accuracv: 0.9996 - loss: 0.0027 - val accuracv: 0.9798 - val loss: 0.0803

#LSTM with Dropout

```

model3 = Sequential([
    Embedding(max_words, 64, input_length=max_len),
    LSTM(64, return_sequences=False),
    Dropout(0.5),
    Dense(1, activation='sigmoid')
])
model3.compile(loss='binary_crossentropy', optimizer='adam', metrics=['accuracy'])
model3.summary()

```

Train

```

model3.fit(X_train_pad, y_train, epochs=5, batch_size=32, validation_split=0.1)

```

Model: "sequential_2"

Layer (type)	Output Shape	Param #
embedding_2 (Embedding)	?	0 (unbuilt)
lstm_2 (LSTM)	?	0 (unbuilt)
dropout (Dropout)	?	0
dense_2 (Dense)	?	0 (unbuilt)

Total params: 0 (0.00 B)

Trainable params: 0 (0.00 B)

Non-trainable params: 0 (0.00 B)

Epoch 1/5

126/126 ————— 13s 74ms/step - accuracy: 0.8552 - loss: 0.4512 - val_accuracy: 0.8565 - val_loss: 0.4149

Epoch 2/5

126/126 ————— 8s 62ms/step - accuracy: 0.8625 - loss: 0.4045 - val_accuracy: 0.8565 - val_loss: 0.4148

Epoch 3/5

126/126 ————— 21s 143ms/step - accuracy: 0.8668 - loss: 0.3968 - val_accuracy: 0.8565 - val_loss: 0.4125

Epoch 4/5

126/126 ————— 7s 58ms/step - accuracy: 0.8558 - loss: 0.4194 - val_accuracy: 0.8565 - val_loss: 0.4145

Epoch 5/5

126/126 ————— 10s 59ms/step - accuracy: 0.8738 - loss: 0.3801 - val_accuracy: 0.8565 - val_loss: 0.4132

```
def evaluate_model(model, X_test_pad, y_test):
    y_pred = (model.predict(X_test_pad) > 0.5).astype("int32")
    print("Accuracy:", accuracy_score(y_test, y_pred))
    print("Precision:", precision_score(y_test, y_pred))
    print("F1 Score:", f1_score(y_test, y_pred))
    print("\nClassification Report:\n", classification_report(y_test, y_pred))


    cm = confusion_matrix(y_test, y_pred)
    plt.figure(figsize=(6,4))
    sns.heatmap(cm, annot=True, fmt="d", cmap="Blues", xticklabels=['Ham', 'Spam'], yticklabels=['Ham', 'Spam'])
    plt.xlabel("Predicted")
    plt.ylabel("Actual")
    plt.title("Confusion Matrix")
    plt.show()

evaluate_model(model1, X_test_pad, y_test)
```

35/35 ————— 2s 40ms/step
Accuracy: 0.8654708520179372
Precision: 0.0
F1 Score: 0.0

Classification Report:				
	precision	recall	f1-score	support
0	0.87	1.00	0.93	965
1	0.00	0.00	0.00	150
accuracy			0.87	1115
macro avg	0.43	0.50	0.46	1115
weighted avg	0.75	0.87	0.80	1115


```
/usr/local/lib/python3.11/dist-packages/sklearn/metrics/_classification.py:1565: UndefinedMetricWarning: Precision is ill-defined ar
_warn_prf(average, modifier, f"{metric.capitalize()} is", len(result))
/usr/local/lib/python3.11/dist-packages/sklearn/metrics/_classification.py:1565: UndefinedMetricWarning: Precision is ill-defined ar
_warn_prf(average, modifier, f"{metric.capitalize()} is", len(result))
/usr/local/lib/python3.11/dist-packages/sklearn/metrics/_classification.py:1565: UndefinedMetricWarning: Precision is ill-defined ar
_warn_prf(average, modifier, f"{metric.capitalize()} is", len(result))
/usr/local/lib/python3.11/dist-packages/sklearn/metrics/_classification.py:1565: UndefinedMetricWarning: Precision is ill-defined ar
_warn_prf(average, modifier, f"{metric.capitalize()} is", len(result))
```



evaluate_model(model2, X_test_pad, y_test)

35/35 ————— 3s 63ms/step
Accuracy: 0.9820627802690582
Precision: 0.9779411764705882
F1 Score: 0.9300699300699301

Classification Report:				
	precision	recall	f1-score	support
0	0.98	1.00	0.99	965
1	0.98	0.89	0.93	150
accuracy			0.98	1115
macro avg	0.98	0.94	0.96	1115
weighted avg	0.98	0.98	0.98	1115



```
evaluate_model(model3, X_test_pad, y_test)
```

35/35 ————— 2s 53ms/step

Accuracy: 0.8654708520179372

Precision: 0.0

F1 Score: 0.0

Classification Report:

	precision	recall	f1-score	support
0	0.87	1.00	0.93	965
1	0.00	0.00	0.00	150
accuracy			0.87	1115
macro avg	0.43	0.50	0.46	1115
weighted avg	0.75	0.87	0.80	1115

/usr/local/lib/python3.11/dist-packages/sklearn/metrics/_classification.py:1565: UndefinedMetricWarning: Precision is ill-defined ar
warn prf(average, modifier, f"{metric.capitalize()} is", len(result))